

ML project

Me

12/08/2021

Reading the given training and test sets

```
knitr::opts_chunk$set(echo = TRUE)
library(caret)
```

```
## Loading required package: lattice
```

```
## Loading required package: ggplot2
```

```
library(ggplot2)

pml_training<- read.csv('pml-training.csv')
pml_testing<- read.csv('pml-testing.csv')
```

Having a quick look at our training and test sets

```
head(pml_training)
str(pml_training)
head(pml_testing)
str(pml_testing)
```

Cleaning the data

Step 1: Dealing with NA values

NA strategy: if a variable or column has more than 50% NA values then we will omit it from our dataset. Any variable that satisfies this condition and has less than or equal to 50% NA values will be kept. After that we will impute the remaining NA values and check if it will be a viable predictor for our model building or not

```
cond<- colSums(is.na(pml_training))<= (0.50*nrow(pml_training))
my_cleanset<- pml_training[, cond] #subsetting our training set to include only those variables that satisfy cond i.e. we will only keep a variable or column that has 50% or lesser NA values
```

We are now checking for any integer or numeric variables that are being stored in character format and convert them into numeric class. We will apply the same condition from before i.e. we will only keep a variable or column that has 50% or lesser NA values

```
character_variables<- my_cleanset[,apply(my_cleanset, class)=='character']
str(character_variables)
```

We will now apply our NA strategy to see which variables to keep and which variables to remove. "" is the same as NA value in character format

```
cond2<- colSums(character_variables=="")<= (0.50*nrow(character_variables)) #we want to find
  how many of the character variables can be converted into numeric class provided they have a
  tleast 50% of the values i.e. we will tolerate maximum 50% NA values
cond2
character_variables_subset<- character_variables[,cond2] #subsetting character_variables so t
  hat it satisfies our cond2 or the NA strategy
head(character_variables_subset) #so we don't need any of the columns aside from user_name, c
  vtd_timestamp, new_window and classe as they did not satisfy cond2
```

Now doing the appropriate class conversions

```
library(dplyr)
```

```
##
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':
##
##   filter, lag
```

```
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```
library(lubridate)
```

```
##
## Attaching package: 'lubridate'
```

```
## The following objects are masked from 'package:base':
##
##   date, intersect, setdiff, union
```

```
character_variables_subset <- mutate(character_variables_subset, cvtd_timestamp= dmy_hm(cvtd_
  timestamp))
str(character_variables_subset)
```

Now replacing the proper preprocessed columns in character_variables_subset to my_cleanset or our training set

```
my_cleanset2<- my_cleanset
my_cleanset2$cvtd_timestamp<- character_variables_subset$cvtd_timestamp
my_cleanset<- my_cleanset2
```

Now we are removing all the unnecessary character variables that did not satisfy cond2

```
char_var_to_be_removed<- names(character_variables[,!cond2])  
my_cleanset<- select(my_cleanset, -char_var_to_be_removed)
```

```
## Note: Using an external vector in selections is ambiguous.  
## i Use `all_of(char_var_to_be_removed)` instead of `char_var_to_be_removed` to silence this  
message.  
## i See <https://tidyselect.r-lib.org/reference/faq-external-vector.html>.  
## This message is displayed once per session.
```

Step 2: Data conversions and removing remaining unnecessary variables**

Converting classe and new_window into factor variables

```
my_cleanset$classe<- as.factor(my_cleanset$classe)  
my_cleanset$new_window<- as.factor(my_cleanset$new_window)
```

We will use `nearZeroVar()` to find out the variables that have near zero variance and will omit them from our model building process

```
library(ISLR)  
library(caret)
```

```
## Loading required package: lattice
```

```
## Loading required package: ggplot2
```

```
nearZeroVar(my_cleanset, saveMetrics = TRUE) #so the variable new_window has near zero variance. So it won't be a good predictor for our model.
```

##	freqRatio	percentUnique	zeroVar	nzv
## X	1.000000	100.00000000	FALSE	FALSE
## user_name	1.100679	0.03057792	FALSE	FALSE
## raw_timestamp_part_1	1.000000	4.26562022	FALSE	FALSE
## raw_timestamp_part_2	1.000000	85.53154622	FALSE	FALSE
## cvtd_timestamp	1.000668	0.10192641	FALSE	FALSE
## new_window	47.330049	0.01019264	FALSE	TRUE
## num_window	1.000000	4.37264295	FALSE	FALSE
## roll_belt	1.101904	6.77810621	FALSE	FALSE
## pitch_belt	1.036082	9.37722964	FALSE	FALSE
## yaw_belt	1.058480	9.97349913	FALSE	FALSE
## total_accel_belt	1.063160	0.14779329	FALSE	FALSE
## gyros_belt_x	1.058651	0.71348486	FALSE	FALSE
## gyros_belt_y	1.144000	0.35164611	FALSE	FALSE
## gyros_belt_z	1.066214	0.86127816	FALSE	FALSE
## accel_belt_x	1.055412	0.83579655	FALSE	FALSE
## accel_belt_y	1.113725	0.72877383	FALSE	FALSE
## accel_belt_z	1.078767	1.52379982	FALSE	FALSE
## magnet_belt_x	1.090141	1.66649679	FALSE	FALSE
## magnet_belt_y	1.099688	1.51870350	FALSE	FALSE
## magnet_belt_z	1.006369	2.32901845	FALSE	FALSE
## roll_arm	52.338462	13.52563449	FALSE	FALSE
## pitch_arm	87.256410	15.73234125	FALSE	FALSE
## yaw_arm	33.029126	14.65701763	FALSE	FALSE
## total_accel_arm	1.024526	0.33635715	FALSE	FALSE
## gyros_arm_x	1.015504	3.27693405	FALSE	FALSE
## gyros_arm_y	1.454369	1.91621649	FALSE	FALSE
## gyros_arm_z	1.110687	1.26388747	FALSE	FALSE
## accel_arm_x	1.017341	3.95984099	FALSE	FALSE
## accel_arm_y	1.140187	2.73672409	FALSE	FALSE
## accel_arm_z	1.128000	4.03628580	FALSE	FALSE
## magnet_arm_x	1.000000	6.82397309	FALSE	FALSE
## magnet_arm_y	1.056818	4.44399144	FALSE	FALSE
## magnet_arm_z	1.036364	6.44684538	FALSE	FALSE
## roll_dumbbell	1.022388	84.20650290	FALSE	FALSE
## pitch_dumbbell	2.277372	81.74498012	FALSE	FALSE
## yaw_dumbbell	1.132231	83.48282540	FALSE	FALSE
## total_accel_dumbbell	1.072634	0.21914178	FALSE	FALSE
## gyros_dumbbell_x	1.003268	1.22821323	FALSE	FALSE
## gyros_dumbbell_y	1.264957	1.41677709	FALSE	FALSE
## gyros_dumbbell_z	1.060100	1.04984201	FALSE	FALSE
## accel_dumbbell_x	1.018018	2.16593619	FALSE	FALSE
## accel_dumbbell_y	1.053061	2.37488533	FALSE	FALSE
## accel_dumbbell_z	1.133333	2.08949139	FALSE	FALSE
## magnet_dumbbell_x	1.098266	5.74864948	FALSE	FALSE
## magnet_dumbbell_y	1.197740	4.30129447	FALSE	FALSE
## magnet_dumbbell_z	1.020833	3.44511263	FALSE	FALSE
## roll_forearm	11.589286	11.08959331	FALSE	FALSE
## pitch_forearm	65.983051	14.85577413	FALSE	FALSE
## yaw_forearm	15.322835	10.14677403	FALSE	FALSE
## total_accel_forearm	1.128928	0.35674243	FALSE	FALSE
## gyros_forearm_x	1.059273	1.51870350	FALSE	FALSE
## gyros_forearm_y	1.036554	3.77637346	FALSE	FALSE
## gyros_forearm_z	1.122917	1.56457038	FALSE	FALSE
## accel_forearm_x	1.126437	4.04647844	FALSE	FALSE
## accel_forearm_y	1.059406	5.11160942	FALSE	FALSE
## accel_forearm_z	1.006250	2.95586586	FALSE	FALSE

```
## magnet_forearm_x      1.012346    7.76679238  FALSE FALSE
## magnet_forearm_y      1.246914    9.54031189  FALSE FALSE
## magnet_forearm_z      1.000000    8.57710733  FALSE FALSE
## classe                 1.469581    0.02548160  FALSE FALSE
```

#We also note that the variable X contains just the serial numbers so we won't need it.

so removing X and new_window from our train set

```
my_cleanset_updated<- select(my_cleanset, -c(X, new_window))
str(my_cleanset_updated) #this is the final cleaned and formatted train set
```

```
## 'data.frame':    19622 obs. of  58 variables:
## $ user_name      : chr  "carlitos" "carlitos" "carlitos" "carlitos" ...
## $ raw_timestamp_part_1: int  1323084231 1323084231 1323084231 1323084232 1323084232 13230
84232 1323084232 1323084232 1323084232 1323084232 ...
## $ raw_timestamp_part_2: int  788290 808298 820366 120339 196328 304277 368296 440390 4843
23 484434 ...
## $ cvtd_timestamp    : POSIXct, format: "2011-12-05 11:23:00" "2011-12-05 11:23:00" ...
## $ num_window        : int  11 11 11 12 12 12 12 12 12 12 ...
## $ roll_belt         : num  1.41 1.41 1.42 1.48 1.48 1.45 1.42 1.42 1.43 1.45 ...
## $ pitch_belt        : num  8.07 8.07 8.07 8.05 8.07 8.06 8.09 8.13 8.16 8.17 ...
## $ yaw_belt          : num  -94.4 -94.4 -94.4 -94.4 -94.4 -94.4 -94.4 -94.4 -94.4 -94.4
...
## $ total_accel_belt   : int  3 3 3 3 3 3 3 3 3 3 ...
## $ gyros_belt_x       : num  0 0.02 0 0.02 0.02 0.02 0.02 0.02 0.02 0.03 ...
## $ gyros_belt_y       : num  0 0 0 0 0.02 0 0 0 0 0 ...
## $ gyros_belt_z       : num  -0.02 -0.02 -0.02 -0.03 -0.02 -0.02 -0.02 -0.02 -0.02 0 ...
## $ accel_belt_x       : int  -21 -22 -20 -22 -21 -21 -22 -22 -20 -21 ...
## $ accel_belt_y       : int  4 4 5 3 2 4 3 4 2 4 ...
## $ accel_belt_z       : int  22 22 23 21 24 21 21 21 24 22 ...
## $ magnet_belt_x      : int  -3 -7 -2 -6 -6 0 -4 -2 1 -3 ...
## $ magnet_belt_y      : int  599 608 600 604 600 603 599 603 602 609 ...
## $ magnet_belt_z      : int  -313 -311 -305 -310 -302 -312 -311 -313 -312 -308 ...
## $ roll_arm           : num  -128 -128 -128 -128 -128 -128 -128 -128 -128 -128 ...
## $ pitch_arm          : num  22.5 22.5 22.5 22.1 22.1 22 21.9 21.8 21.7 21.6 ...
## $ yaw_arm            : num  -161 -161 -161 -161 -161 -161 -161 -161 -161 -161 ...
## $ total_accel_arm    : int  34 34 34 34 34 34 34 34 34 34 ...
## $ gyros_arm_x        : num  0 0.02 0.02 0.02 0 0.02 0 0.02 0.02 0.02 ...
## $ gyros_arm_y        : num  0 -0.02 -0.02 -0.03 -0.03 -0.03 -0.03 -0.02 -0.03 -0.03 ...
## $ gyros_arm_z        : num  -0.02 -0.02 -0.02 0.02 0 0 0 0 -0.02 -0.02 ...
## $ accel_arm_x        : int  -288 -290 -289 -289 -289 -289 -289 -289 -288 -288 ...
## $ accel_arm_y        : int  109 110 110 111 111 111 111 111 109 110 ...
## $ accel_arm_z        : int  -123 -125 -126 -123 -123 -122 -125 -124 -122 -124 ...
## $ magnet_arm_x       : int  -368 -369 -368 -372 -374 -369 -373 -372 -369 -376 ...
## $ magnet_arm_y       : int  337 337 344 344 337 342 336 338 341 334 ...
## $ magnet_arm_z       : int  516 513 513 512 506 513 509 510 518 516 ...
## $ roll_dumbbell      : num  13.1 13.1 12.9 13.4 13.4 ...
## $ pitch_dumbbell     : num  -70.5 -70.6 -70.3 -70.4 -70.4 ...
## $ yaw_dumbbell       : num  -84.9 -84.7 -85.1 -84.9 -84.9 ...
## $ total_accel_dumbbell: int  37 37 37 37 37 37 37 37 37 37 ...
## $ gyros_dumbbell_x   : num  0 0 0 0 0 0 0 0 0 0 ...
## $ gyros_dumbbell_y   : num  -0.02 -0.02 -0.02 -0.02 -0.02 -0.02 -0.02 -0.02 -0.02 -0.02
...
## $ gyros_dumbbell_z   : num  0 0 0 -0.02 0 0 0 0 0 0 ...
## $ accel_dumbbell_x   : int  -234 -233 -232 -232 -233 -234 -232 -234 -232 -235 ...
## $ accel_dumbbell_y   : int  47 47 46 48 48 48 47 46 47 48 ...
## $ accel_dumbbell_z   : int  -271 -269 -270 -269 -270 -269 -270 -272 -269 -270 ...
## $ magnet_dumbbell_x  : int  -559 -555 -561 -552 -554 -558 -551 -555 -549 -558 ...
## $ magnet_dumbbell_y  : int  293 296 298 303 292 294 295 300 292 291 ...
## $ magnet_dumbbell_z  : num  -65 -64 -63 -60 -68 -66 -70 -74 -65 -69 ...
## $ roll_forearm       : num  28.4 28.3 28.3 28.1 28 27.9 27.9 27.8 27.7 27.7 ...
## $ pitch_forearm      : num  -63.9 -63.9 -63.9 -63.9 -63.9 -63.9 -63.9 -63.8 -63.8 -63.8
...
## $ yaw_forearm        : num  -153 -153 -152 -152 -152 -152 -152 -152 -152 -152 ...
## $ total_accel_forearm: int  36 36 36 36 36 36 36 36 36 36 ...
## $ gyros_forearm_x    : num  0.03 0.02 0.03 0.02 0.02 0.02 0.02 0.02 0.03 0.02 ...
## $ gyros_forearm_y    : num  0 0 -0.02 -0.02 0 -0.02 0 -0.02 0 0 ...
## $ gyros_forearm_z    : num  -0.02 -0.02 0 0 -0.02 -0.03 -0.02 0 -0.02 -0.02 ...
```

```
## $ accel_forearm_x      : int  192 192 196 189 189 193 195 193 193 190 ...
## $ accel_forearm_y      : int  203 203 204 206 206 203 205 205 204 205 ...
## $ accel_forearm_z      : int  -215 -216 -213 -214 -214 -215 -215 -213 -214 -215 ...
## $ magnet_forearm_x     : int  -17 -18 -18 -16 -17 -9 -18 -9 -16 -22 ...
## $ magnet_forearm_y     : num   654 661 658 658 655 660 659 660 653 656 ...
## $ magnet_forearm_z     : num   476 473 469 469 473 478 470 474 476 473 ...
## $ classe               : Factor w/ 5 levels "A","B","C","D",...: 1 1 1 1 1 1 1 1 1 1 ...
```

Building a model with classe as our outcome variable and random forest as our prediction method

Algorithm: We will use random forest here as our outcome variable classe is a categorical variable and random forests are good with non linear data

Now dividing the train set further into a smaller train (70%) and validation set(30%). We will then building a model with random forest method

```
set.seed(100)
project_train<- createDataPartition(my_cleaset_updated$classe, p=0.70, list = FALSE)
my_cleaset_training<- my_cleaset_updated[project_train,]
my_cleaset_validation<- my_cleaset_updated[-project_train,]
```

Using parallel processing for improving the processing time of random forest

```
#Step 1: Configure parallel processing
library(parallel)
library(doParallel)
```

```
## Loading required package: foreach
```

```
## Loading required package: iterators
```

```
cluster <- makeCluster(detectCores() - 1) # convention to leave 1 core for OS
registerDoParallel(cluster)

#Configuring trainControl object.We will be doing cross validation with 10 folds

modControl<- trainControl(method = 'cv', number = 5, verboseIter = TRUE, allowParallel = TRUE
)

#Finally, building a model with random forest method

set.seed(100)
system.time(model_RF<- train(classe~., data= my_cleaset_training, method= 'rf', trControl= modControl))
```

```
## Aggregating results
## Selecting tuning parameters
## Fitting mtry = 31 on full training set
```

```
##      user  system elapsed
## 39.51    1.74  576.51
```

#Step 4: De-register parallel processing cluster

```
stopCluster(cluster)
registerDoSEQ()
```

Now evaluating the accuracy on our validation set

```
set.seed(100)
modelRF_predictions<- predict(model_RF, my_cleaset_validation)
confusionMatrix(my_cleaset_validation$classe, modelRF_predictions)
```

```
## Confusion Matrix and Statistics
##
##              Reference
## Prediction      A      B      C      D      E
##      A 1674      0      0      0      0
##      B      0 1139      0      0      0
##      C      0      2 1024      0      0
##      D      0      0      1  963      0
##      E      0      0      0      0 1082
##
## Overall Statistics
##
##              Accuracy : 0.9995
##              95% CI : (0.9985, 0.9999)
##      No Information Rate : 0.2845
##      P-Value [Acc > NIR] : < 2.2e-16
##
##              Kappa : 0.9994
##
##      McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##              Class: A Class: B Class: C Class: D Class: E
## Sensitivity          1.0000  0.9982  0.9990  1.0000  1.0000
## Specificity          1.0000  1.0000  0.9996  0.9998  1.0000
## Pos Pred Value       1.0000  1.0000  0.9981  0.9990  1.0000
## Neg Pred Value       1.0000  0.9996  0.9998  1.0000  1.0000
## Prevalence           0.2845  0.1939  0.1742  0.1636  0.1839
## Detection Rate       0.2845  0.1935  0.1740  0.1636  0.1839
## Detection Prevalence 0.2845  0.1935  0.1743  0.1638  0.1839
## Balanced Accuracy     1.0000  0.9991  0.9993  0.9999  1.0000
```

Conclusion

So we are getting an accuracy of 0.9985 or near 100% accuracy (approximately). So we can conclude that using a random forest model is giving us near perfect accuracy on our validation set

Prediciting on the test data set now

```
pml_testing_updated<- pml_testing[,-160] #removing problem ID from test set
#cleaning the test set
pml_testing_updated<- select(pml_testing_updated, names(my_cleanset_updated)[-58]) #selecting
only the variables in our final cleaned train set and removing classe as it does not exist th
e test set

pml_testing_updated<- mutate(pml_testing_updated, cvtd_timestamp= dmy_hm(cvtd_timestamp))
set.seed(100)
#now predicting on the test set
testing_predictions<- predict(model_RF, pml_testing_updated)
testing_predictions
```

```
## [1] B A B A A E D B A A B C B A E E A B B B
## Levels: A B C D E
```